

# Advanced Digital Systems Design with FPGAs


## EECS 395/495 – Lecture 8

Dr. David Zaretsky  
[david.zaretsky@northwestern.edu](mailto:david.zaretsky@northwestern.edu)


# Computer Networking

- There are various models of the networking stack, typically arranged in 3-7 layers.


OSI 7-layer model



DOD 3-layer model



Simplified 4/5-layer model



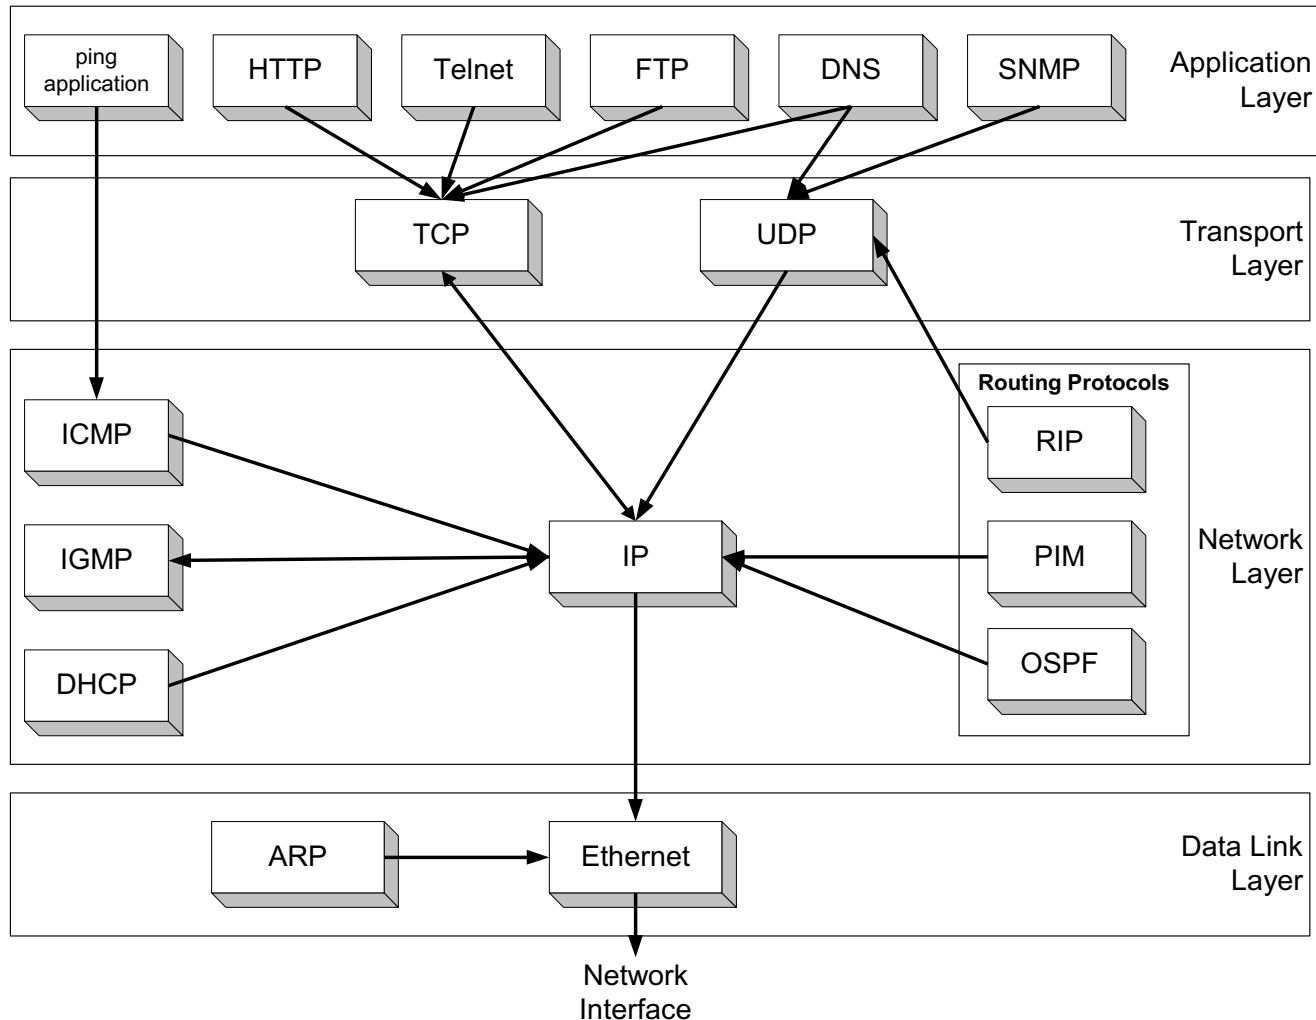


# TCP/IP Network Model

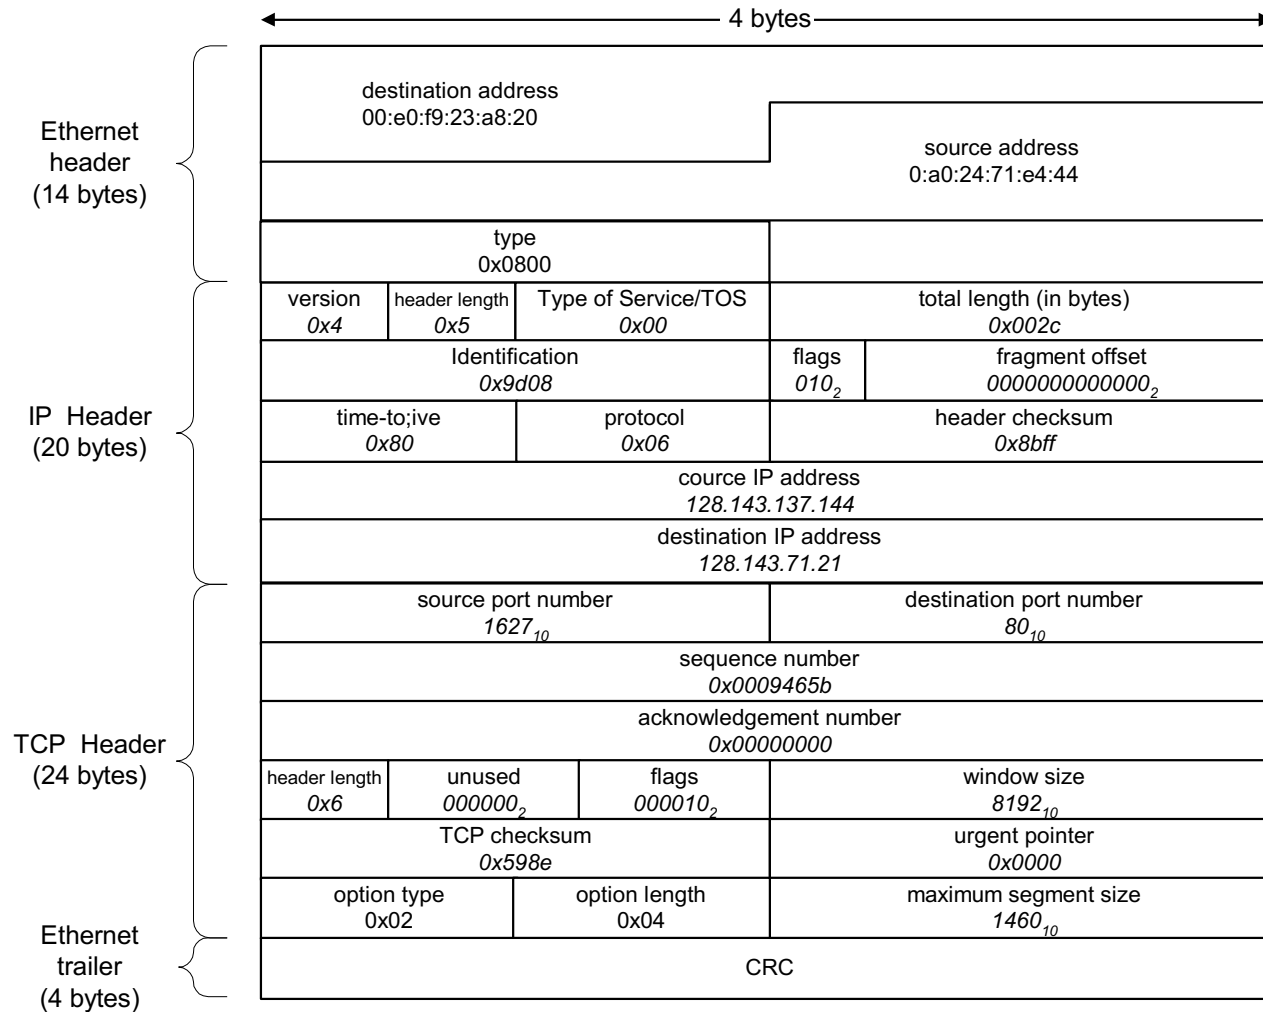
---

- ▶ Physical Layer – the physical wires and hardware
- ▶ Link Layer - includes device driver and network interface card
- ▶ Network Layer - handles the movement of packets, i.e. Routing
- ▶ Transport Layer - provides a reliable flow of data between two hosts
- ▶ Application Layer - handles the details of the particular application

# Assignment of Protocols to Layers

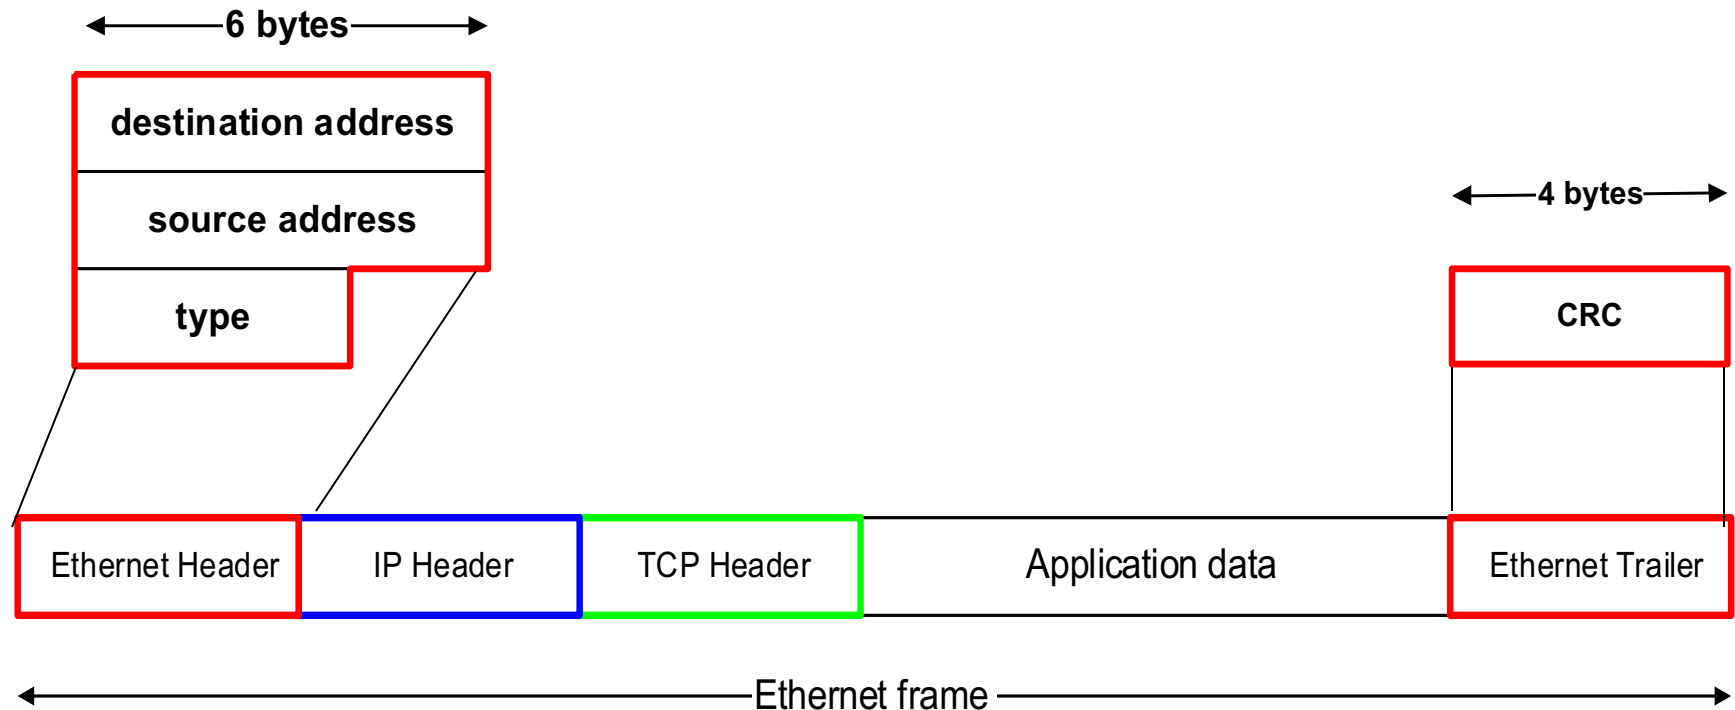


# Parsing the Packets



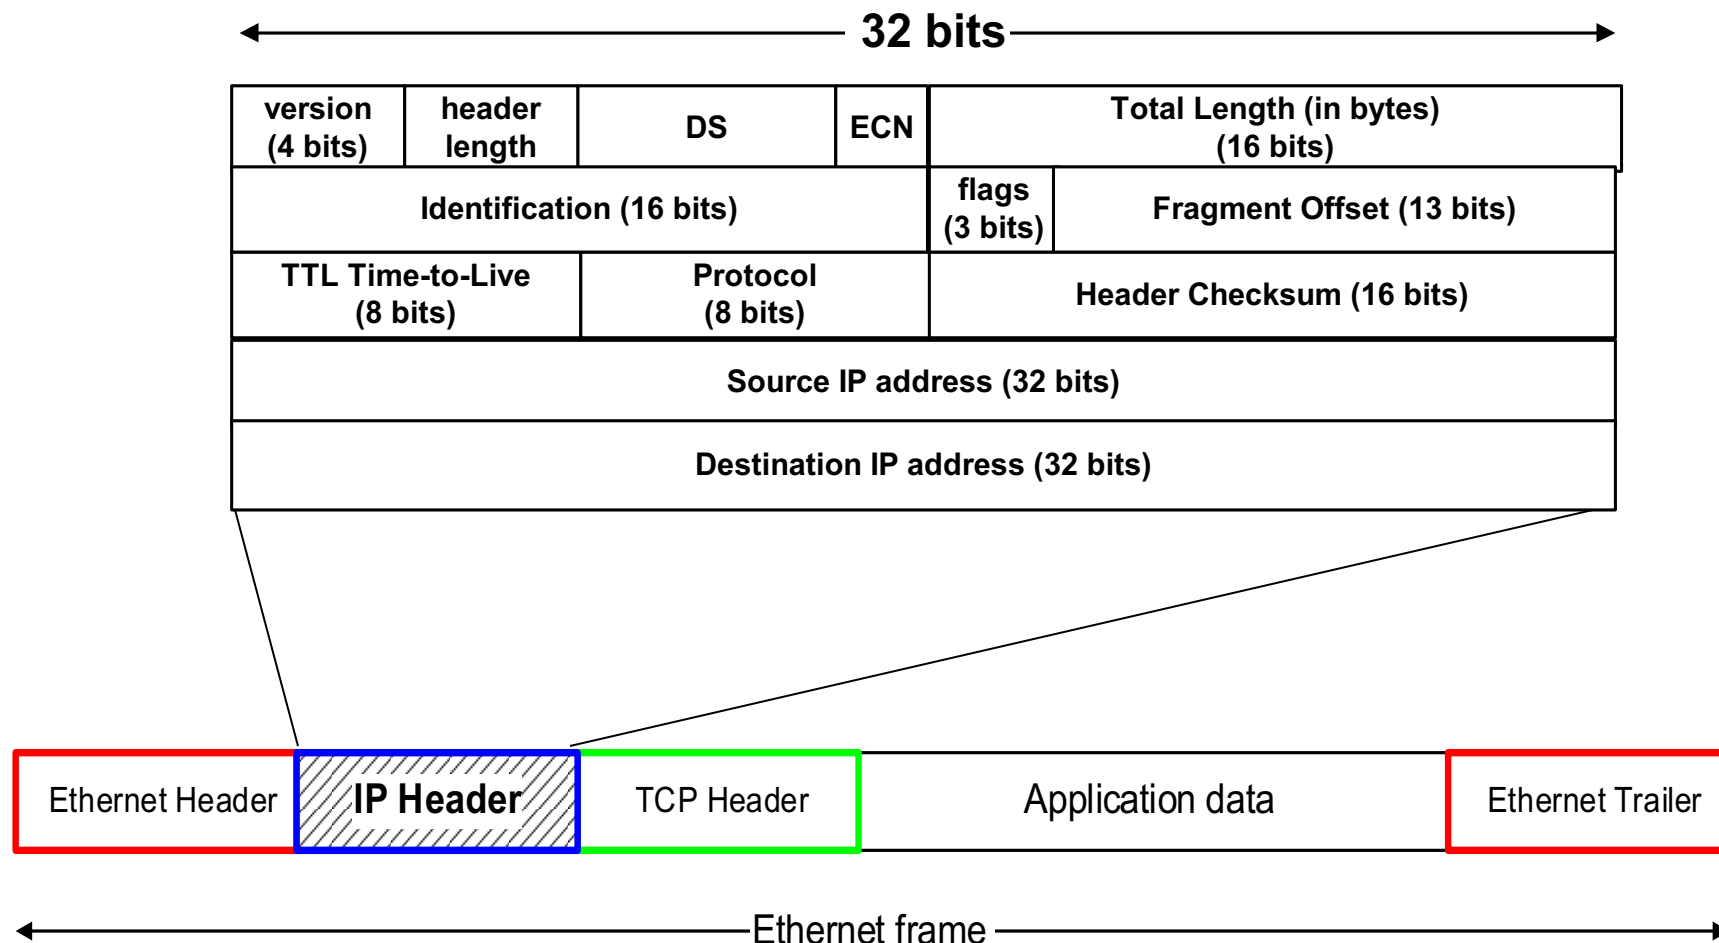
# Ethernet Header

- ▶ 14 byte header
- ▶ 4 byte checksum



# IP Header

## ► 20 byte header




# TCP Header

← 32 bits →

Source Port Number		Destination Port Number	
Sequence number (32 bits)			
Acknowledgement number (32 bits)			
header length	0	Flags	window size
TCP checksum			urgent pointer
option type	length		Max. segment size

**Option:**  
maximum  
segment size



← Ethernet frame →





# Internet Protocol (IP)

---

- ▶ IP is responsible for addressing and routing of data packets.
- ▶ Two versions in current in use: IPv4 & IPv6.
- ▶ IPv4: uses a 160 bit (20 byte) header, and 32 bit addresses.
- ▶ IPv6 was mainly developed to increase IP address space due to the huge growth in Internet usage during the 1990s.
- ▶ IPv6 uses a 320 bit (40 byte) header and 128 bit addresses.
- ▶ Header fields include: source and destination addresses, packet length and packet number.



# IP Header Fields

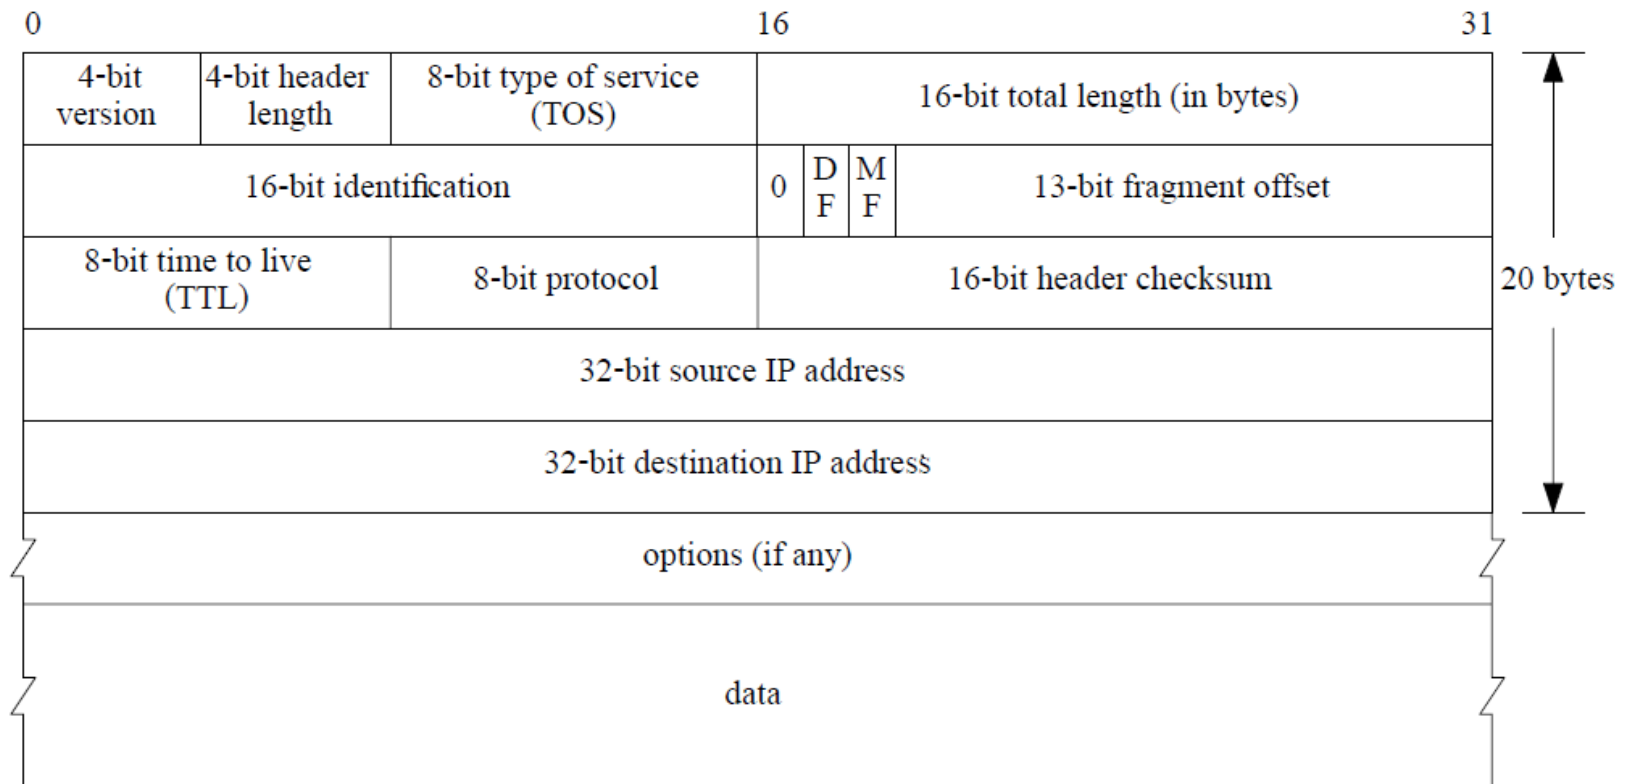
---

- ▶ Ver – version of IP
- ▶ IHL – Internet Header Length (32-bit words)
- ▶ Service – Precedence/Delay/Throughput/Reliability
- ▶ Identification – assistance in reassembling fragments
- ▶ CF – control flags:
  - ▶ Reserved
  - ▶ I to prevent fragmentation, else 0
  - ▶ I if last fragment, else 0
- ▶ Fragment Offset – of this fragment in total message, bytes
- ▶ TTL – Time to Live, upper limit of life enroute
- ▶ Protocol – next higher protocol, e.g., TCP, UDP or ICMP



# IP Datagram

## IP Header





# Transport Control Protocol (TCP)

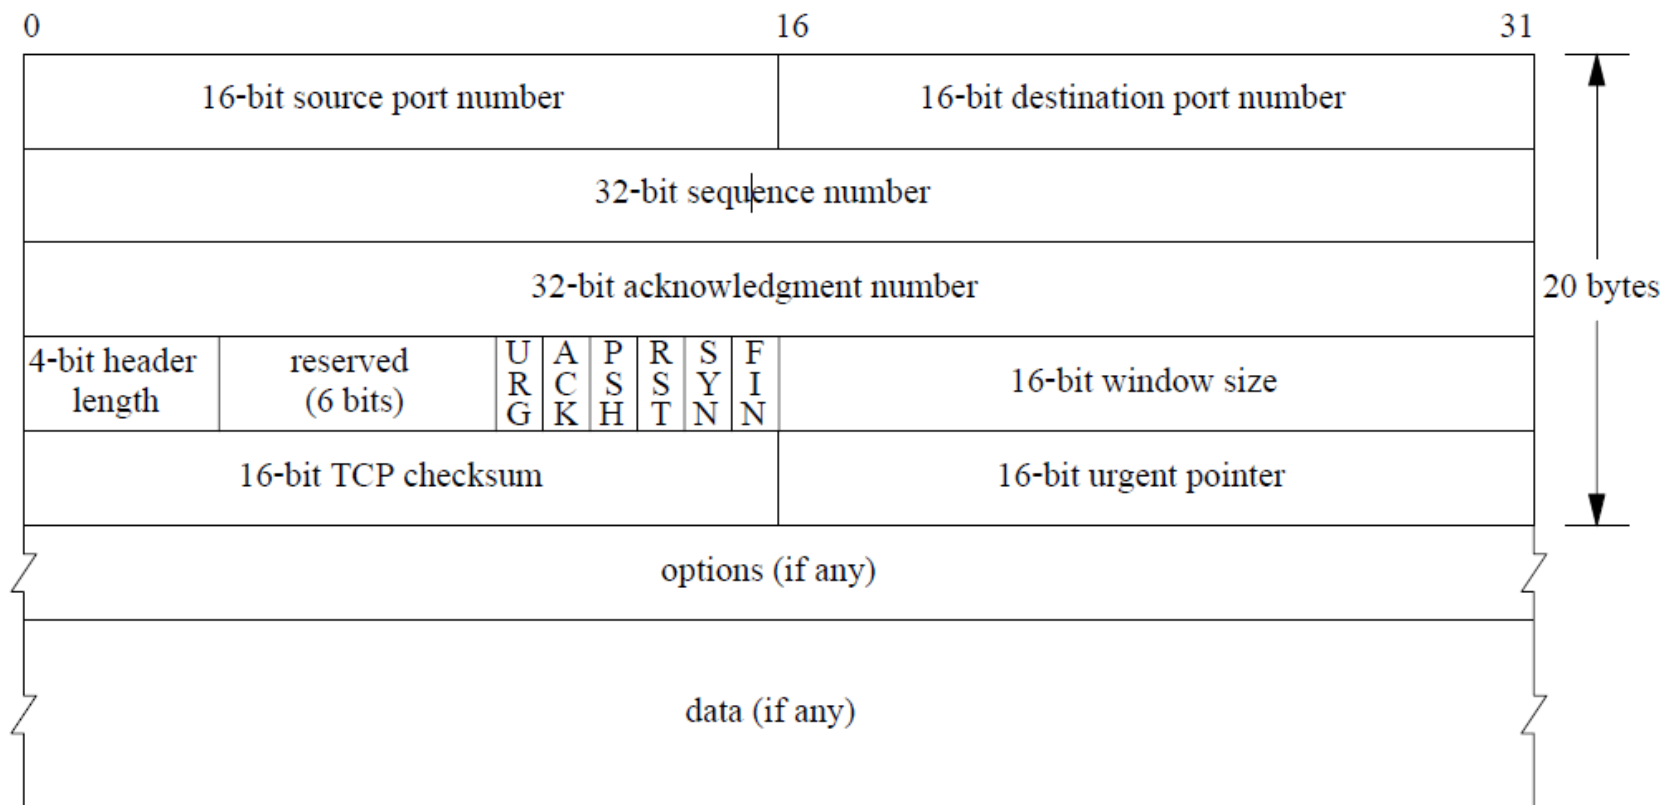
---

- ▶ Reliable, full-duplex, connection-oriented, stream delivery
- ▶ Data is guaranteed to arrive, and in the correct order without duplications
- ▶ Imposes significant overheads
- ▶ Connections are established using a three-way handshake
- ▶ Data is divided up into packets by the operating system
- ▶ Packets are numbered, and received packets are acknowledged
- ▶ Connections are explicitly closed



# TCP Datagram

## TCP Header

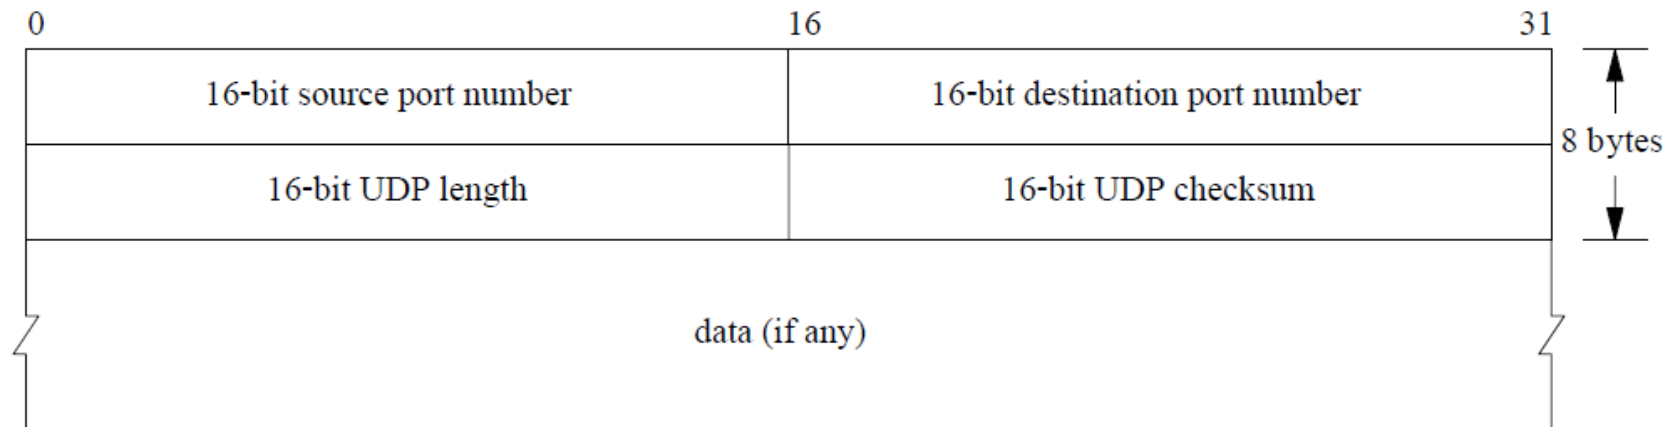




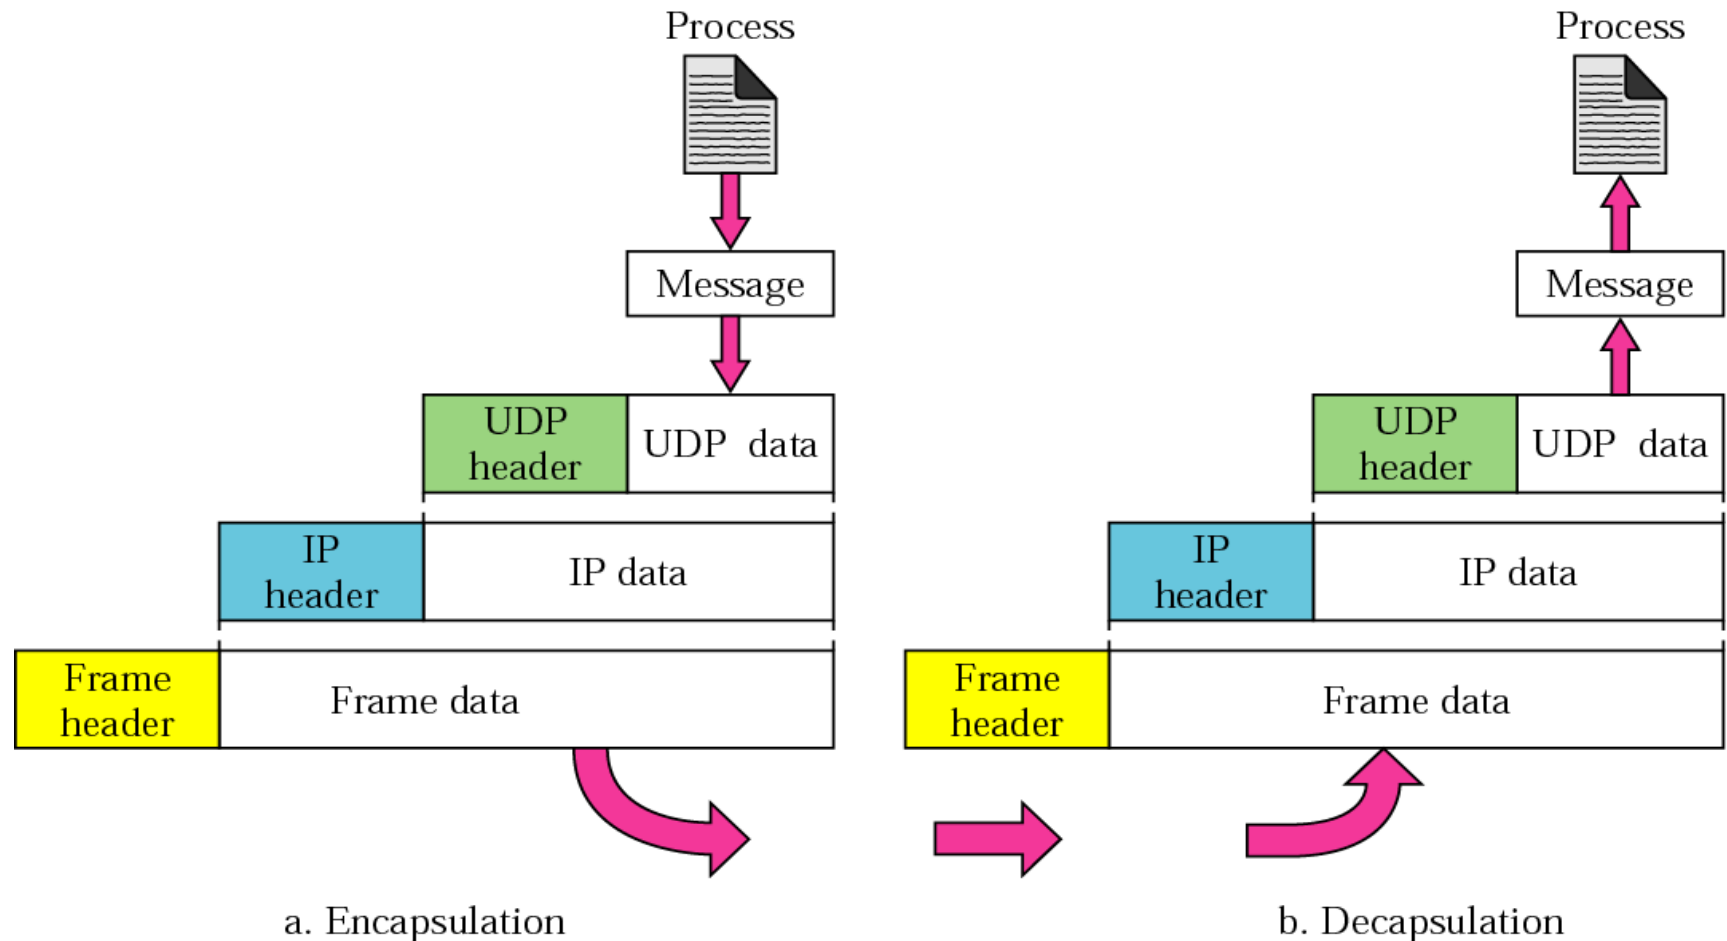
# User Datagram Protocol (UDP)

---

- ▶ One-to-one or one-to-many, connectionless and unreliable protocol
- ▶ Adds packet length + checksum to guard against corrupted packets
- ▶ Source and destination ports are used to associate a packet with a specific application at each end
- ▶ Not guaranteed to arrive, in order, or lossless
- ▶ Use Cases
  - ▶ Where packet loss is better handled by the application than the network stack
  - ▶ Where the overhead of setting up a connection isn't wanted
- ▶ Typical Applications
  - ▶ VOIP
  - ▶ Audio / video simulcasting

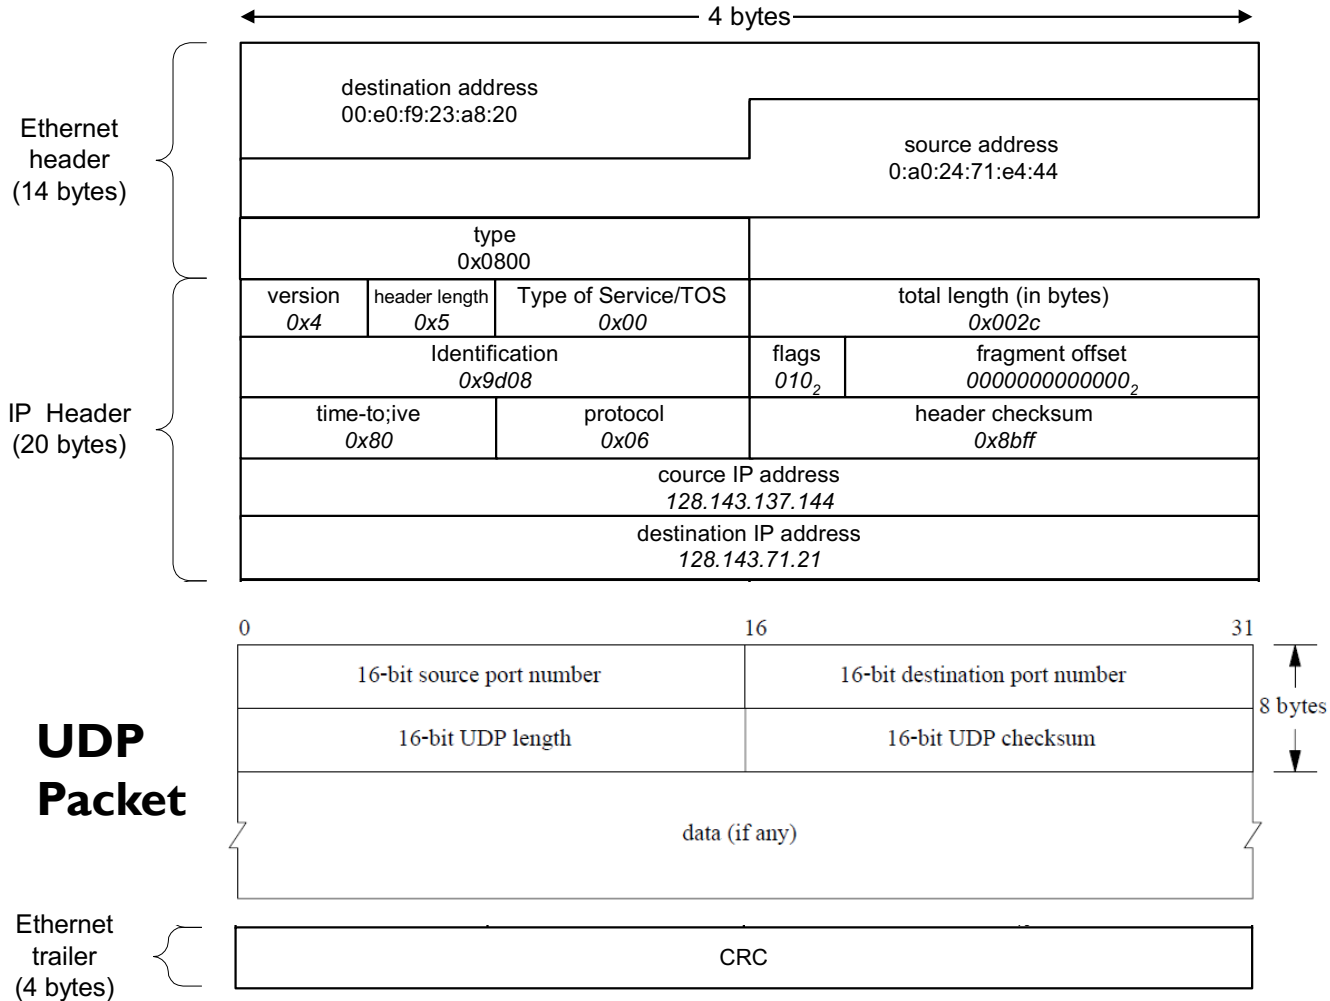


# UDP Encapsulation / Decapsulation





# Parsing UDP Packets





# UDP Reader in C

```
#define ETH_DST_ADDR_BYTES 6
#define ETH_SRC_ADDR_BYTES 6
#define ETH_PROTOCOL_BYTES 2
#define IP_VERSION_BYTES 1
#define IP_HEADER_BYTES 1
#define IP_TYPE_BYTES 1
#define IP_LENGTH_BYTES 2
#define IP_ID_BYTES 2
#define IP_FLAG_BYTES 2
#define IP_TIME_BYTES 1
#define IP_PROTOCOL_BYTES 1
#define IP_CHECKSUM_BYTES 2
#define IP_SRC_ADDR_BYTES 4
#define IP_DST_ADDR_BYTES 4
#define UDP_DST_PORT_BYTES 2
#define UDP_SRC_PORT_BYTES 2
#define UDP_LENGTH_BYTES 2
#define UDP_CHECKSUM_BYTES 2
#define IP_PROTOCOL_DEF 0x0800
#define IP_VERSION_DEF 0x4
#define IP_HEADER_LENGTH_DEF 0x5
#define IP_TYPE_DEF 0x0
#define IP_FLAGS_DEF 0x4
#define TIME_TO_LIVE 0xe
#define UDP_PROTOCOL_DEF 0x11

int read_udp_packet(FILE *source, unsigned char *packet_data)
{
    unsigned char eth_dst_addr[ETH_DST_ADDR_BYTES];
    unsigned char eth_src_addr[ETH_SRC_ADDR_BYTES];
    unsigned char eth_protocol[ETH_PROTOCOL_BYTES];
    unsigned char ip_version[IP_VERSION_BYTES];
    unsigned char ip_header[IP_HEADER_BYTES];
    unsigned char ip_type[IP_TYPE_BYTES];
    unsigned char ip_length[IP_LENGTH_BYTES];
    unsigned char ip_id[IP_ID_BYTES];

    unsigned char ip_flag[IP_FLAG_BYTES];
    unsigned char ip_time[IP_TIME_BYTES];
    unsigned char ip_protocol[IP_PROTOCOL_BYTES];
    unsigned char ip_checksum[IP_CHECKSUM_BYTES];
    unsigned char ip_dst_addr[IP_SRC_ADDR_BYTES];
    unsigned char ip_src_addr[IP_DST_ADDR_BYTES];
    unsigned char udp_dst_port[UDP_DST_PORT_BYTES];
    unsigned char udp_src_port[UDP_SRC_PORT_BYTES];
    unsigned char udp_length[UDP_LENGTH_BYTES];
    unsigned char udp_checksum[UDP_CHECKSUM_BYTES];
    unsigned char udp_data[1024];
    unsigned short udp_data_length = 0, crc = 0, checksum = 0;
    int p = 0;

    if ( feof(source) ) return 0;

    fread(eth_dst_addr, 1, ETH_DST_ADDR_BYTES, source);
    fread(eth_src_addr, 1, ETH_SRC_ADDR_BYTES, source);
    fread(eth_protocol, 1, ETH_PROTOCOL_BYTES, source);
    if ( (((unsigned int)eth_protocol[0] << 8) | (unsigned
int)eth_protocol[1]) != IP_PROTOCOL_DEF )
        return 0;

    fread(ip_version, 1, IP_VERSION_BYTES, source);
    if ( (ip_version[0] >> 4) != IP_VERSION_DEF )
        return 0;
    ip_header[0] = ip_version[0] & 0xF;

    fread(ip_type, 1, IP_TYPE_BYTES, source);
    fread(ip_length, 1, IP_LENGTH_BYTES, source);
    fread(ip_id, 1, IP_ID_BYTES, source);
    fread(ip_flag, 1, IP_FLAG_BYTES, source);
    fread(ip_time, 1, IP_TIME_BYTES, source);
    fread(ip_protocol, 1, IP_PROTOCOL_BYTES, source);
    if ( ip_protocol[0] != UDP_PROTOCOL_DEF )
        return 0;

    fread(ip_checksum, 1, IP_CHECKSUM_BYTES, source);
    fread(ip_src_addr, 1, IP_SRC_ADDR_BYTES, source);
    fread(ip_dst_addr, 1, IP_DST_ADDR_BYTES, source);
    fread(udp_dst_port, 1, UDP_DST_PORT_BYTES, source);
    fread(udp_src_port, 1, UDP_SRC_PORT_BYTES, source);
    fread(udp_length, 1, UDP_LENGTH_BYTES, source);
    fread(udp_checksum, 1, UDP_CHECKSUM_BYTES, source);

    // get the UDP data
    udp_data_length = (((unsigned int)udp_length[0] << 8) |
(unsigned int)udp_length[1]);
    udp_data_length -= (UDP_CHECKSUM_BYTES +
UDP_LENGTH_BYTES + UDP_DST_PORT_BYTES +
UDP_SRC_PORT_BYTES); // + CME_HEADER);
    fread(udp_data, 1, udp_data_length, source);

    // calculate the checksum
    crc = udp_sum_calc( ip_src_addr, ip_dst_addr, ip_protocol,
ip_length, udp_src_port, udp_dst_port, udp_length, udp_data
);
    checksum = ((unsigned int)udp_checksum[0] << 8) |
(unsigned int)udp_checksum[1];
    if ( checksum != crc ) {
        fprintf( stderr, "ERROR: Checksum mismatch -- %04x !=
%04x\n", crc, checksum);
        return 0;
    }

    for ( int i = 0; i < udp_data_length; i++ ) {
        packet_data[i] = udp_data[i];
    }

    return udp_data_length;
}
```





# UDP Checksum

```
unsigned short udp_sum_calc(
    unsigned char *ip_src_addr,
    unsigned char *ip_dst_addr,
    unsigned char *ip_protocol,
    unsigned char *ip_length,
    unsigned char *udp_src_port,
    unsigned char *udp_dst_port,
    unsigned char *udp_length,
    unsigned char *udp_data)
{
    unsigned short padd = 0;
    unsigned int sum = 0;
    int udp_len = ((udp_length[0] << 8) | udp_length[1]);
    int data_length = (UDP_CHECKSUM_BYTES +
        UDP_LENGTH_BYTES + UDP_DST_PORT_BYTES +
        UDP_SRC_PORT_BYTES);

    // Find out if the length of data is even or odd number. If odd,
    // add a padding byte = 0 at the end of packet
    if ( (data_length & 1) == 1 ) {
        padd=1;
        udp_data[data_length]=0;
    }

    // add the UDP pseudo header
    for (int i=0; i<4; i=i+2) {
        sum += (((ip_src_addr[i]<<8)&0xFF00)+(ip_src_addr[i+1]&0xFF));
    }

    for (int i=0; i<4; i=i+2) {
        sum += (((ip_dst_addr[i]<<8)&0xFF00)+(ip_dst_addr[i+1]&0xFF));
    }
}
```

```
sum += ip_protocol[0];
sum += (((unsigned short)ip_length[0]<<8)&0xFF00)+(ip_length[1]&0xFF) - 20;
sum += (((unsigned short)udp_dst_port[0]<<8)&0xFF00) +
(udp_dst_port[1]&0xFF);
sum += (((unsigned short)udp_src_port[0]<<8)&0xFF00) +
(udp_src_port[1]&0xFF);
sum += (((unsigned short)udp_length[0]<<8)&0xFF00) +
(udp_length[1]&0xFF);

// make 16 bit words out of every two adjacent 8 bit words and
// calculate the sum of all 16 bit words
for (int i=0; i < data_length; i += 2)
{
    sum += (((unsigned short)udp_data[i]<<8)&0xFF00);
    sum += (udp_data[i+1]&0xFF);
}

// keep only the last 16 bits of the 32 bit calculated sum and add the carries
while (sum >> 16 != 0)
{
    sum = (sum & 0xFFFF) + (sum >> 16);
}

// Take the one's complement of sum
sum = ~sum;
//printf("sum: %08x\n", sum );

return sum;
}
```

# Programming Assignment #3: Streaming UDP Reader

---



- ▶ **Streaming UDP process**
  - ▶ Read packets from an input FIFO
  - ▶ Extract datagrams and write to output FIFO
  - ▶ Use 0x02 and 0x03 values to delineate start and end of packets, respectively.
- ▶ **Simulate & Synthesize**
  - ▶ Use pcap file as input (will need to remove pcap header)
  - ▶ Simulate to get cycle count and verify correctness
  - ▶ Synthesize to get resource utilization

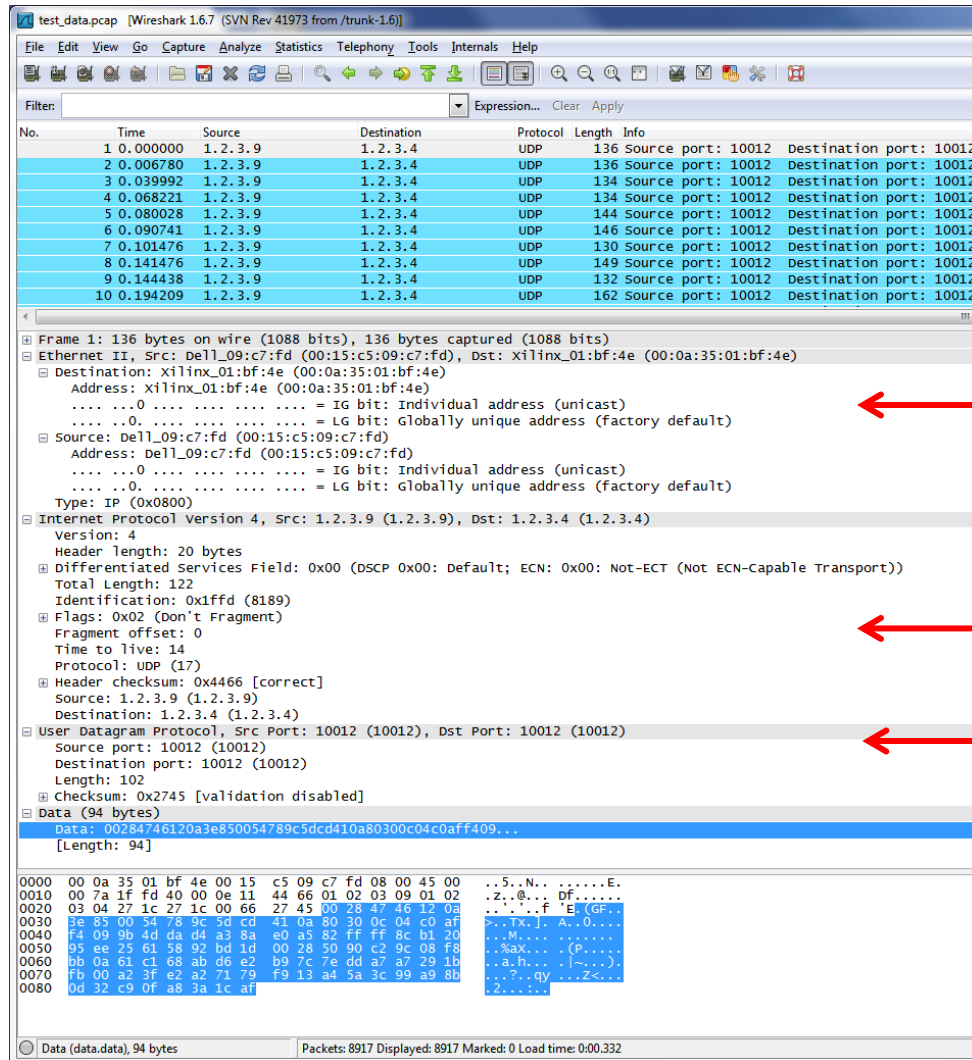


# Implementation of UDP Reader

---

- ▶ Synchronize packets using start of frame (0x02) and end of frame (0x03) delineators
- ▶ Checksum can be calculated in parallel as each data point is acquired, instead of doing it at the end
- ▶ Data needs to be validated before it goes out
  - ▶ Store in temporary fifo buffer
  - ▶ Clear the fifo if any checksum errors are found
  - ▶ Burst out packets after checksum is validated
- ▶ Use the REPORT command in VHDL to display data in the log window

# Wireshark



test\_data.pcap [Wireshark 1.6.7 (SVN Rev 41973 from /trunk-1.6)]

File Edit View Go Capture Analyze Statistics Telephony Tools Internals Help

Filter: Expression... Clear Apply

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000	1.2.3.9	1.2.3.4	UDP	136	Source port: 10012 Destination port: 10012
2	0.006780	1.2.3.9	1.2.3.4	UDP	136	Source port: 10012 Destination port: 10012
3	0.039992	1.2.3.9	1.2.3.4	UDP	134	Source port: 10012 Destination port: 10012
4	0.068221	1.2.3.9	1.2.3.4	UDP	134	Source port: 10012 Destination port: 10012
5	0.080028	1.2.3.9	1.2.3.4	UDP	144	Source port: 10012 Destination port: 10012
6	0.090741	1.2.3.9	1.2.3.4	UDP	146	Source port: 10012 Destination port: 10012
7	0.101476	1.2.3.9	1.2.3.4	UDP	130	Source port: 10012 Destination port: 10012
8	0.141476	1.2.3.9	1.2.3.4	UDP	149	Source port: 10012 Destination port: 10012
9	0.144438	1.2.3.9	1.2.3.4	UDP	132	Source port: 10012 Destination port: 10012
10	0.194209	1.2.3.9	1.2.3.4	UDP	162	Source port: 10012 Destination port: 10012

Frame 1: 136 bytes on wire (1088 bits), 136 bytes captured (1088 bits)

Ethernet II, Src: Dell\_09:c7:fd (00:15:c5:09:c7:fd), Dst: xilinx\_01:bf:4e (00:0a:35:01:bf:4e)

Destination: xilinx\_01:bf:4e (00:0a:35:01:bf:4e)

Address: xilinx\_01:bf:4e (00:0a:35:01:bf:4e)

...0... = IG bit: Individual address (unicast)

...0... = LG bit: Globally unique address (factory default)

Source: Dell\_09:c7:fd (00:15:c5:09:c7:fd)

Address: Dell\_09:c7:fd (00:15:c5:09:c7:fd)

...0... = IG bit: Individual address (unicast)

...0... = LG bit: Globally unique address (factory default)

Type: IP (0x0800)

Internet Protocol Version 4, Src: 1.2.3.9 (1.2.3.9), Dst: 1.2.3.4 (1.2.3.4)

Version: 4

Header length: 20 bytes

Differentiated Services Field: 0x00 (DSCP 0x00: Default; ECN: 0x00: Not-ECT (Not ECN-Capable Transport))

Total Length: 122

Identification: 0x1ffd (8189)

Flags: 0x02 (Don't Fragment)

Fragment offset: 0

Time to live: 14

Protocol: UDP (17)

Header checksum: 0x4466 [correct]

Source: 1.2.3.9 (1.2.3.9)

Destination: 1.2.3.4 (1.2.3.4)

User Datagram Protocol, Src Port: 10012 (10012), Dst Port: 10012 (10012)

Source port: 10012 (10012)

Destination port: 10012 (10012)

Length: 102

Checksum: 0x2745 [validation disabled]

Data (94 bytes)

Data: 00284746120a3e850054789c5dc410a80300c04c0aff409...

[Length: 94]

0000 00 0a 35 01 bf 4e 00 15 c5 09 c7 fd 08 00 45 00 ..S..N.....E.

0010 00 7a 1f fd 40 00 0e 11 44 66 01 02 03 09 01 02 .Z.@..DF.....

0020 03 04 27 1c 27 1c 00 66 27 45 00 28 47 46 12 0a ...fE(GF.....

0030 3e 85 00 54 78 9c 5d cd 41 0a 80 30 0c 04 c0 af >..TX..J..A..0...

0040 f4 09 9b 4d da d4 a3 8a e0 a5 82 ff ff 8c b1 20 ..M.....

0050 95 ee 23 61 58 92 bd 1d 00 28 30 90 c2 9c 08 f8 ...ax... (P....

0060 bb 0a 61 c1 68 ab d6 e2 b9 7c 7e dd a7 a7 29 1b ...a.h...[~....

0070 fb 00 a2 3f e2 a2 71 79 f9 13 a4 5a 3c 99 a9 8b ...?..qy...Z<...

0080 0d 32 c9 0f a8 3a 1c af 2.....

Data (data.data), 94 bytes Packets: 8917 Displayed: 8917 Marked: 0 Load time: 0:00.332

← Ethernet Frame

← IP Frame

← UDP Frame



# UDP Datagram Output

```
00 28 47 46 12 0a 3e 85 00 54 78 9c 5d cd 41 0a 00 28 47 49 12 0a 3e ca 00 52 78 9c 65 8d 41 0a
80 30 0c 04 c0 af f4 09 9b 4d da d4 a3 8a e0 a5 80 40 0c 03 bf e2 13 d2 b4 da f6 a8 b2 e0 65 2f
82 ff ff 8c b1 20 95 ee 25 61 58 92 bd 1d 00 28 fe ff 31 ae 0b a2 8b b9 84 0c 03 d9 6b 01 22 08
50 90 c2 9c 08 f8 bb 0a 61 c1 68 ab d6 e2 b9 7c 59 92 a2 74 02 0e 52 84 30 0a 50 57 4d 18 e6 87
7e dd a7 a7 29 1b fb 00 a2 3f e2 a2 71 79 f9 13 97 93 87 4f 63 36 f6 6a f6 37 d2 77 8c 30 d8 be
a4 5a 3c 99 a9 8b 0d 32 c9 0f a8 3a 1c af 7e 48 6f dd 5e 64 79 01 bc a3 1c d2

00 28 47 47 12 0a 3e 8c 00 54 78 9c 5d cc bb 0d 00 28 47 4a 12 0a 3e d5 00 5c 78 9c 65 8d c1 0d
80 30 10 83 e1 55 32 82 ed dc e5 92 12 10 12 4d 80 30 0c 03 57 e9 08 b6 53 9a 7e 01 f5 d9 0f fb
2a f6 df 85 87 14 21 f2 37 96 be c2 5b df 71 47 0f 43 41 20 35 aa 3f 56 2e 89 7d f6 06 d4 2a d0
64 77 51 9e 04 04 24 52 30 11 e8 8b 85 1a 31 7c 5d 34 73 01 0e 89 14 b2 08 f4 9d e0 56 ec e7 ed
3f 8f 48 53 ab de 79 6e be 68 16 40 ad c0 0f 0b 52 f7 14 75 e8 35 8e f3 59 23 02 c8 91 55 db 0c
1c 33 31 1c 2c 65 48 ae d1 2e a5 c3 1c c5 0b 2a 8f 4f 9c 48 69 9e e2 f6 7b d2 9a e3 88 8d
d4 0d bc 37 2a 56

00 28 47 48 12 0a 3e ae 00 52 78 9c 5d 8b c1 0d 00 28 47 4b 12 0a 3e e0 00 5e 78 9c 65 8c 3b 0e
80 20 14 43 57 61 84 b6 f2 f9 72 54 43 e2 85 93 c0 20 0c 43 af c2 11 6c f3 09 6b 5b 31 b2 f4 fe
fb ef 22 e0 41 e5 25 4d d3 a6 3d 6a 41 83 58 cc 87 29 a9 54 09 8a 87 58 7a 79 c9 d5 1b 50 ab 40
44 e5 20 c0 21 91 42 14 81 ba 45 57 ee f1 e9 cb 33 31 26 13 60 90 48 21 89 40 3f 08 e6 92 3e de
75 7a 98 d8 35 8c f8 61 6a 9f ae 0f 8c 69 5a 8d 6e 75 0b 6b 4e bd c5 a1 cf 61 19 23 af ac c6 1c
4a d6 96 e9 7d da 7a 03 a2 d1 1c a0 b1 21 17 31 71 22 84 c9 f0 b5 b6 a3 1f 71 64 de
69 fa a3 07 c1 3a 2a 61
```



# UDP Parser in VHDL

```
when WAIT_FOR_SOF_STATE =>
  if ( input_empty = '0' ) then
    input_rd_en <= '1';
  end if;
  if ( (input_empty = '0') and (input_dout = X"02") ) then
    next_state <= ETH_DST_ADDR_STATE;
  end if;

when ETH_DST_ADDR_STATE =>
  if ( input_empty = '0' ) then
    input_rd_en <= '1';
    eth_dst_addr_t := std_logic_vector((unsigned(eth_dst_addr) sll 8) or resize(unsigned(input_dout),ETH_DST_ADDR_BYTES*8));
    eth_dst_addr_c := eth_dst_addr_t;
    num_bytes_c <= (num_bytes + 1) mod ETH_DST_ADDR_BYTES;
    if ( num_bytes = ETH_DST_ADDR_BYTES-1 ) then
      next_state <= ETH_SRC_ADDR_STATE;
    end if;
  end if;

:

when IP_SRC_ADDR_STATE =>
  if ( input_empty = '0' ) then
    input_rd_en <= '1';
    ip_src_addr_t := std_logic_vector((unsigned(ip_src_addr) sll 8) or resize(unsigned(input_dout),IP_SRC_ADDR_BYTES*8));
    ip_src_addr_c <= ip_src_addr_t;
    num_bytes_c <= (num_bytes + 1) mod IP_SRC_ADDR_BYTES;
    if ( num_bytes = IP_SRC_ADDR_BYTES-1 ) then
      checksum_c <= checksum + std_logic_vector(resize(unsigned(ip_src_addr_t(31 downto 16)),32) + resize(unsigned(ip_src_addr_t(15 downto 0)),32));
      next_state <= IP_DST_ADDR_STATE;
    end if;
  end if;
```